

SMMP: A Secure Multi-Party Messaging Protocol

David R. Andersen* and Mark S. Andersland

Department of Electrical and Computer Engineering,

The University of Iowa, Iowa City, IA 52242

Tycho J. Andersen

Canonical, Inc.

Madison, WI 53703

(Dated: March 25, 2014)

Abstract

We describe a new secure, multi-party, synchronous communication protocol. The protocol follows a peer-to-peer model and provides perfect forward secrecy and perfect future secrecy against a computationally-bounded adversary, as well as information-theoretic plausible deniability for participants in a multi-party conversation. The protocol uses a three-round, authenticated Burmester-Desmedt group key agreement protocol to generate a shared secret between a group of N participants. Conversation participants authenticate to each other during key agreement via a triple Diffie-Hellman mechanism. Individual message keys are updated after receipt of each message by incorporating new key material distributed by the sending participant. No security requirements are imposed on the underlying transport layer, and our protocol leaks no metadata beyond that exposed by the transport layer. Conversation transcript universality is assured by a conversation digest that is updated upon receipt of each message, and all conversation messages are signed to verify proof of origin. All setup operations prior to group key agreement take place over an insecure channel. SMMP represents a significant improvement in security and efficiency over current secure, multi-party messaging protocols including GOTR, mpOTR, and improved GOTR.

I. INTRODUCTION

Secure multi-party messaging has been an elusive goal of security researchers for a long time. The initial group off-the-record (GOTR) ‘virtual server’ approach¹ was deemed unsatisfactory because it has the unfortunate drawback of not permitting all participants to confirm that they are receiving unmodified copies of all messages. A further attempt to solve this problem² was incomplete because of the lack of a simple, secure key-agreement strategy. Work on the problem continues to date, most notably with³. Liu *et.al.*⁴ recently improved the security of GOTR/mpOTR using a Burmester-Desmedt group key agreement mechanism imposed on top of a network comprised of binary encrypted channels between participants.

In this paper, we describe a novel secure multi-party messaging protocol (SMMP) with a simple key-agreement algorithm that provides perfect forward secrecy (PFS) and perfect future secrecy (PFuS) against a computationally-bounded adversary, as well as information-theoretic plausible deniability (PD) for participants in a group conversation. Our protocol is a peer-to-peer protocol, whereby each participant (peer) contributes new key material to the group key set with each message sent. This key-update mechanism is called a ratchet. In contrast to the various mpOTR/GOTR proposals, no constraints on the security of the underlying transport layer are required during key agreement or otherwise.

SMMP makes use of six underlying cryptographic primitives, namely: 1) an underlying (unspecified) symmetric encryption algorithm, 2) the elliptic-curve Diffie-Hellman key agreement protocol, 3) the finite-field Diffie-Hellman key agreement protocol, 4) a secure cryptographic hash function, 5) a secure cryptographic message authentication code, and 6) a secure key-derivation function.

II. KEY POINTS

- Burmester-Desmedt⁵ group key agreement takes place using finite-field, cyclic-group Diffie-Hellman key agreement as the underlying protocol.
- Participants authenticate each other during group key agreement via a triple Diffie-Hellman mechanism.
- During the conversation, elliptic curve Diffie-Hellman key operations are performed,

minimizing the required key-exchange bandwidth.

- The protocol is synchronous. A method for resynchronizing those participants that lose synchronization due to lost packets from collisions or transport failure is included in the protocol.
- Group setup is via an insecure channel.
- No security constraints are imposed on the underlying transport layer.
- Participants may be added to the group via a reinitialization of the group secret. A participant may leave the group by simply stopping participation and not updating his ratchet state.
- A uniform message transcript is assured over all participants via a conversation digest that is updated with each message. If the digest does not match what a receiver expects, a resynchronization of the protocol is requested.
- The underlying symmetric encryption algorithm used for communication is unspecified.

III. BACKGROUND

Various protocols exist for facilitating secure one-to-one messaging, including Off the Record Messaging (OTR)^{8–12}, Silent Circle Instant Messaging Protocol (SCIMP)¹³, and Axolotl¹⁴. Each of these protocols provides security through the use of ephemeral keys for symmetric encryption. The method of generating and exchanging these keys varies widely between the three protocols. All of the protocols provide forward secrecy. Some also provide plausible deniability, and one (Axolotl) also provides perfect future secrecy (compromise of the current keyset does not cause compromise of future keysets). However, each of these protocols is limited to the one-to-one messaging case.

OTR was the first such protocol. It is based on an Advertise \rightarrow Acknowledge \rightarrow Use method for key updating. OTR has been widely used for secure instant messaging and applications containing OTR plugins are available for XMPP, IRC, and other messaging systems.

SCIMP is a proprietary protocol developed by Silent Circle as their solution for providing secure communications as part of their product line. The key advancement algorithm for SCIMP is essentially a hash. Each symmetric encryption key is hashed to obtain the key for the next ratchet step. As a result, a key compromise at any point will permit an attacker to follow the conversation from that point on.

Axolotl is the first one-to-one messaging protocol to provide future secrecy. In this protocol, randomly-generated key data is mixed in to the ratchet state with each message send/receive operation, so that a compromised keyset permits an attacker access to only one message. Following the conversation requires a new key compromise with each message sent.

The first proposal for a secure group ($N > 2$) messaging protocol was that of Bian *et.al.*¹ [Group Off The Record (GOTR)]. The protocol is based on a virtual server, a participant that sets up a secure one-on-one channel with each member of the group. When members of the group who are not the virtual server wish to communicate with each other, they relay their messages through the virtual server. Several problems with this approach were identified. First, the virtual server must be assumed honest. The virtual server handles decryption/reencryption from one user to another, as well as authentication of all group members. As a result, a dishonest server (or compromise of the server) would be catastrophic to the security of the group conversation. Second, the protocol leaks metadata. High-level addressing data is sent in plaintext. And finally, there is no mechanism for conversation transcript verification. As a result of these problems with GOTR, a better solution was necessary.

A second proposal for a secure group messaging protocol was made by Goldberg *et.al.*² [multi-party Off The Record (mpOTR)]. While this protocol solves some of the problems associated with the Bian messaging scheme, it has additional problems and is not a good solution. In particular, mpOTR provides PD against offline judges only. Further, it links all messages in a given transcript together, meaning that if an honest user commits to authoring or participation in a part of the transcript, they have committed to the entire transcript. Also, as with all OTR-based protocols, PD is achieved by publishing ephemeral keys post-conversation. Finally, the group encryption key is not automatically updated in mpOTR, and as a result, PFS is not guaranteed - a “window of compromise” exists.

An improvement of the security of the mpOTR/GOTR protocol was made by Liu *et.al.*⁴

using Burmester-Desmedt⁵ for group key agreement. The improved mpOTR/GOTR requires an underlying network of secure, binary communication channels that exist between participants. Complexity of this improved mpOTR/GOTR protocol is an issue. For each group key update, a total of six rounds of group key agreement exchanges are required. Further, each participant must maintain the status for $2(N - 1)$ virtual nodes. Thus, group key agreement is a computationally-intensive task that may preclude conversation participants from updating their encryption keys on a per-message basis. Finally, as with all OTR-based schemes, publication of the ephemeral signing keys at the end of the conversation is required to facilitate PD.

In this paper, we present a secure, peer-to-peer multi-party messaging protocol (SMMP) that provides PFS, PFuS, and PD, and has a simple key agreement algorithm requiring three exchanges to complete key agreement (including authentication between all peers - see Fig. 1). The protocol allows all group members to confirm that they are receiving a complete transcript of the group conversation, and provides a robust mechanism for resynchronization of the ratchet state if messages are lost (*e.g.* because of failure of the underlying transport mechanism).

IV. NOTATION

The following notation is used:

- N is the total number of participants in the group (including the Organizer)
- \oplus is the bitwise XOR operator
- $\{N_i\}$ is the set of all N_i and $\{N_i\}_m = N_m$.
- $\hat{\mathcal{A}}_i$ is the ratchet state for participant P_i .
- $||$ is the concatenation operator
- $||_i x_i$ means concatenation of all x_i running from the smallest to the largest index i
- Lower case keys are private ECDH (or DH in the case of the group key-agreement handshake) keys.

- Upper case keys are public ECDH (or DH in the case of the group key-agreement handshake) keys.
- $\text{hash}()$ is a secure, one-way cryptographic hash function.
- $\text{mac}(k, m)$ is a secure message authentication code that authenticates a message m using key k .
- $\text{KDF}()$ is a secure key derivation function, *e.g.* pbkdf2.
- $c = e(k, m)$ and $m = d(k, c)$ are the underlying (unspecified) symmetric encryption and decryption algorithms that use a key k to transform a plaintext m into ciphertext c and *vice versa*.
- MK is a master key from which the initial ratchet state $\hat{\mathcal{A}}_i$ is computed by each participant P_i . This parameter is securely erased after computing the initial ratchet state.
- X is the generator for the elliptic curve and $\text{ECDH}()$ is the Diffie-Hellman operator on the elliptic curve, *i.e.*, for private key u , the corresponding public key is $U = \text{ECDH}(u, X)$. The shared secret between y and z is then $\text{ECDH}(y, Z) = \text{ECDH}(z, Y) = \text{ECDH}(y, \text{ECDH}(z, X))$.

V. PROTOCOL STATE

Each Participant P_j will maintain the following state variables in persistent storage:

- RK - the root key
- mk - the message key
- v - the group private key
- conv_digest - the digest of all conversation messages so far
- r_j - the participant private ratchet key (for participant P_j , $\text{ECDH}(r_j, X) = R_j$)
- r_j^{init} - the initial value of participant P_j 's private ratchet key

- $\{R_i\}$ - the set of public ratchet keys from each participant P_i
- $\{R_i^{init}\}$ - the set of initial values of the public ratchet keys from each participant P_i
- j - Participant P_j 's participant index number
- N - the group size
- `group_name` - the group name (this may be different for each participant)
- `resync_required` - a flag used to determine if a resynchronization of the ratchet state is required

VI. KEY AGREEMENT

When a group of participants desires to establish a secure group conversation channel, the following preliminary steps are completed (each participant will be known by a long-term private identity key b_i and the corresponding public identity key B_i - proof of knowledge of the private key constitutes proof of identity):

1. The total number of participants in the group N is determined and this number is distributed to all participants.
2. A unique participant index i is assigned to each participant. The method of assigning this index is application-dependent, and may *e.g.* be assigned based on the ordering of the participant public identity keys.

Following the preliminary steps listed above, each member of the group takes part in an authenticated Burmester-Desmedt group key exchange protocol (see Fig. 2). The authentication step prevents attackers from mounting man-in-the-middle attacks and is a variant of Just and Vaudenay¹⁵. An initial ratchet state \hat{A} is computed from the BD master key.

Each group participant P_j completes the following steps:

1. Participant P_j generates an ephemeral ECDH ratchet key pair (r_j, R_j) , and an ephemeral finite-field DH handshake key pair (h_j, H_j) .
2. Participant P_j broadcasts public identity key B_j and ephemeral public ratchet key R_j .

3. Participant P_j computes the set of keys:

$$k_{TDHi} = \text{hash}(\text{ECDH}(b_j, R_i) \parallel \text{ECDH}(r_j, B_i) \parallel \text{ECDH}(r_j, R_i))$$

and macs:

$$\text{mac}_i = \text{mac}(k_{TDHi}, H_j).$$

4. Participant P_j sends $H_j \parallel \text{mac}_i$ to participant P_i for all participants.

5. Participant P_j receives $H_i \parallel \text{mac}_j$ from each participant P_i .

6. For each participant $i \neq j$, participant P_j computes the set of received keys:

$$k_{r-TDHi} = \text{hash}(\text{ECDH}(r_j, B_i) \parallel \text{ECDH}(b_j, R_i) \parallel \text{ECDH}(r_j, R_i))$$

and macs:

$$\text{mac}_{r-i} = \text{mac}(k_{r-TDHi}, H_i).$$

7. Participant P_j tests if $\text{mac}_{r-i} = \text{mac}_j$. If this test fails, the identity of participant P_i is not confirmed and P_j goes no further.

8. Participant P_j computes: $K_j = (H_{j+1}/H_{j-1})^{h_j} \mod p$ where the index j is taken in a cycle. P_j broadcasts K_j to all participants P_i .

9. Participant P_j receives K_i from each participant P_i and computes:

$$\text{MK} = \text{hash}(H_{j-1}^{Nh_j} \cdot K_j^{N-1} \cdot K_{j+1}^{N-2} \cdots K_{j-2} \mod p),$$

where the index j is taken in a cycle.

10. Participant P_j computes his/her initial ratchet state $\hat{\mathcal{A}}_j$:

$$\text{RK} = \text{KDF}(\text{MK}, 0\text{x}00),$$

$$\text{mk} = \text{KDF}(\text{MK}, 0\text{x}01),$$

$$v = \text{KDF}(\text{MK}, 0\text{x}02),$$

$$\text{conv_digest} = 0\text{x}00 * 32,$$

r_j from step 1,

$$r_j^{\text{init}} = r_j,$$

$\{R_i\}$ from participants,

$$\{R_i^{\text{init}}\} = \{R_i\},$$

$\text{resync_required} = \text{False}.$

VII. SENDING MESSAGES

Participants wanting to send a message proceed as follows:

1. When participant P_j wishes to communicate with other participants, he forms a message m .
2. Participant P_j generates a new ephemeral ratchet key (r_j^{new}, R_j^{new}) .
3. For all $i \neq j$, participant P_j computes the set of $N - 1$ preliminary ciphertexts $c'_i = e(\text{mk}, j \parallel R_j^{new} \oplus \text{hash}(\text{ECDH}(r_j, R_i)) \parallel m)$.
4. Participant P_j computes the set of $N - 1$ macs $c_{maci} = \text{mac}(v, c'_i)$.
5. Participant P_j forms the set of $N - 1$ ciphertexts $c_i = c'_i \parallel c_{maci}$.
6. Participant P_j updates $\text{conv_digest} = \text{hash}(m) \oplus \text{conv_digest}$.
7. Participant P_j updates his/her ratchet state $\hat{\mathcal{A}}_j$ as follows:

$$(r_i, R_i) = (r_i^{new}, R_i^{new}),$$

$$\text{RK} = \text{hash}(\text{RK} \parallel \text{ECDH}(v, \text{conv_digest} \parallel \parallel_i R_i)),$$

$$\text{mk} = \text{KDF}(\text{RK}, 0x01).$$
8. Participant P_j sends c_i to participant P_i for all participants, using a collision-avoidance algorithm that should be transport-specific and is unspecified here.

VIII. RECEIVING MESSAGES

Participants receiving a message proceed as follows:

1. Upon receipt of a ciphertext c , participant P_j separates c into c' and c_{mac} parts.
2. Participant P_j tests if $\text{mac}(v, c') = c_{mac}$. If they do not match, he/she raises a `Bad_HMAC` exception and goes no further.
3. Participant P_j obtains $q \parallel R \parallel m = d(\text{mk}, c')$. If this operations fails, he/she raises a `Message_Undecryptable` exception, sets the `resync_required` flag `True`, and passes c' and control to the message-received housekeeping routine described in Section IX A.

4. Participant P_j sets $R_q = R \oplus \text{hash}(\text{ECDH}(r_j, R_q))$.
5. Participant P_j updates $\text{conv_digest} = \text{hash}(m) \oplus \text{conv_digest}$.
6. Participant P_j updates his/her ratchet state $\hat{\mathcal{A}}_i$ as follows:
 $\text{RK} = \text{hash}(\text{RK} || \text{ECDH}(v, \text{conv_digest} || ||_i R_i))$,
 $\text{mk} = \text{KDF}(\text{RK}, 0x01)$.

IX. PROTOCOL HOUSEKEEPING

Here we list several protocol housekeeping operations that may be necessary. The list may be extended if other useful operations are identified.

A. Message Received Housekeeping

Housekeeping messages will be routed based on a one-byte header prepended to the payload of the housekeeping message. The byte values and their corresponding housekeeping tasks are (currently there are 2):

Byte Value	Housekeeping Task	Location
0x00	Resynchronizing the Ratchet State	Section IX A 1
0x01	Instant Messaging Within SMMP	Section IX A 4

To determine the proper routing, participant P_j proceeds as follows:

1. Participant P_j computes $b || m = d(v, c')$. He/she then finds the value corresponding to byte b in the table above, and sends message m and control to that section. If this decryption fails, participant P_j raises a `Bad_DIGEST` exception and goes no further.

1. *Resynchronizing the Ratchet State*: $b = 0x00$

It is possible, during communication within the group, that a participant's ratchet state may become unsynchronized. Transport layer failures due to lost messages or message collisions can cause a participant to be unsynchronized. If this is the case, a `Message_Undecryptable` exception will be raised and control of the decryption will be passed here. The receiving participant P_j will then proceed as follows:

1. Participant P_j decomposes $m = q || v^{new} || R_q^{new}$.
2. Participant P_j sets $R_q^{init} = R_q^{new}$ and $\{R_i\} = \{R_i^{init}\}$.
3. Participant P_j sets $v = \text{hash}(v || v^{new})$.
4. Participant P_j updates his/her ratchet state $\hat{\mathcal{A}}_j$ as follows:

$$\begin{aligned} \text{RK} &= \text{hash}(v || \text{ECDH}(v, ||_i R_i)), \\ \text{mk} &= \text{KDF}(\text{RK}, 0x01), \\ \text{conv_digest} &= 0x00 * 32, \\ \text{resync_required} &= \text{False}. \end{aligned}$$

2. Adding Members To the Group

Members may be added to the group by running the key agreement part of the protocol again, and including the new participant.

3. Removing Members From the Group

Members may drop from the group simply by not updating their ratchet state. Further, the group may remove a participant by running the key agreement part of the protocol again without the dropped participant.

4. One-To-One Messaging Within SMMP: $b = 0x01$

Participants in the group may exchange one-to-one messages with other group participants using the group infrastructure. If participant P_i wishes to message participant P_j , he/she proceeds as follows:

1. Participant P_i forms a message m .
2. If this is the first private message with P_j , participant P_i makes a copy of his current private ratchet key r_i (s_i) as well as participant P_j 's current public ratchet key R_j (S_j). If this is not the first private message P_i has exchanged with P_j , s_i and S_j already exist.

3. Participant P_i generates a new one-to-one ratchet key (t_i, T_i) .
4. Participant P_i forms preliminary ciphertexts $c' = e(\text{hash}(\text{ECDH}(s_i, S_j), 0x01 \parallel T_i \parallel m))$.
5. Participant P_i computes $c_{mac} = \text{mac}(v, c')$.
6. Participant P_i sends $c = c' \parallel c_{mac}$ to P_j .

Decryption is a straightforward reversal of this process.

B. Message Send Housekeeping

It may be necessary at some point to send housekeeping messages. This section describes procedures to be followed in this case.

1. *Resynchronizing the Ratchet State: resync_required = True*

When the `resync_required` flag is set to *True*, the ratchet state is unsynchronized, and decrypting further conversation messages is impossible. A participant P_j wishing to correct this situation should follow the following procedures (this may be automated):

1. Participant P_j generates a new group private key v^{new} .
2. Participant P_j generates a new ratchet key pair (r^{new}, R^{new}) .
3. participant P_j computes preliminary ciphertext $c' = e(v, 0x00 \parallel j \parallel v^{new} \parallel R^{new})$.
4. Participant P_j computes $\text{mac } c_{mac} = \text{mac}(v, c')$.
5. Participant P_j forms $c = c' \parallel c_{mac}$.
6. Participant P_j waits according to a collision-avoidance algorithm that should be transport-specific and is unspecified here.
7. Participant P_j tests `resync_required` to see if it is still *True*. If not, a resynchronization message was received and no further action is necessary. If *True*, participant P_j continues with the next step.
8. Participant P_j broadcasts c to all participants P_i .

9. Participant P_j updates his/her ratchet state $\hat{\mathcal{A}}_j$ as follows:

$$\begin{aligned}
v &= \text{hash}(v \parallel v^{new}), \\
r_j &= r^{new}, \\
r_j^{init} &= r^{new}, \\
R_j^{init} &= R^{new}, \\
\{R_i\} &= \{R_i^{init}\}, \\
\text{RK} &= \text{hash}(v \parallel \text{ECDH}(v, \parallel_i R_i)), \\
\text{mk} &= \text{KDF}(\text{RK}, 0x01), \\
\text{conv_digest} &= 0x00 * 32, \\
\text{resync_required} &= \text{False}.
\end{aligned}$$

X. SECURITY AND EFFICIENCY ANALYSIS

A. Security Model

Following the work of Goldberg, *et.al.*², we propose a security model with the following attributes:

- Confidentiality - While a participant is willing to disclose certain information to members of a group, the group communication should remain hidden to those outside the group.
- Entity authentication - Prior to joining the group, members should be authenticated so that during the group conversation, group members can be confident that messages purportedly from a particular group member were actually authored by that group member.
- Origin authentication - All messages should be authenticated as to their participant of origin.
- Forward secrecy - The protocol should provide PFS for all messages sent as part of the group communication.
- Future secrecy - The protocol should provide PFuS for all future messages to be sent as part of the group communication.

- **Plausible Deniability** - The protocol should provide PD such that a transcript of the group conversation cannot be used to prove the membership or participation in the group of any participant during the conversation or after the conversation has completed.

SMMP, as described above, achieves each of these requirements.

Further, we adopt a generalized version of the threat model from Goldberg *et.al*² that includes the possibility of both online and offline judges. In particular, we analyze the robustness of the protocol to three types of adversaries, a security adversary, a consensus adversary, and a privacy adversary. Online judges (members of the group) can be considered a subset of the dishonest participants. Both online and offline judges will be used to determine if the goals of the protocol have been met. In the following, we provide a performance analysis of the protocol in the context of these three adversaries.

Security Adversary - The security adversary is a global adversary that is capable of monitoring all network traffic between group participants during the group conversation. The security adversary is not permitted to interact with conversation participants in any way prior to the completion of the conversation. However, following the conversation, he may ask for static or private information of any (or all) group participant(s). The goal of the security adversary is to read messages he is not entitled to, and the security adversary wins if he can read at least one such message.

In our protocol, message privacy is obtained via the use of an unspecified symmetric encryption algorithm where the encryption key changes in an indeterminate way with every message. The security adversary described above may have access to some plaintext as well as all ciphertext messages. Any recovery of private keys following the completion of the group conversation will not provide access to the message encryption keys due to the forward secrecy properties of the protocol. Assuming that the security adversary is computationally bounded and the symmetric encryption algorithm is resistant to known-plaintext attacks, the security adversary will not be able to decrypt and read additional conversation messages beyond those recovered directly from a participant.

Consensus Adversary - The consensus adversary may participate in the conversation or control others in their participation. The goal of the consensus adversary is to get an honest participant to believe that his transcript of the conversation is a match to the group conversation when it does not match.

Our protocol contains a conversation digest mechanism that is updated with the transmission of each conversation message. This digest allows the message receiver to be confident that the message received was identical to the message sent. Further, the digest is updated with each message during the conversation until a resynchronization event is requested. Assuming that the cryptographic hash function used to compute the conversation digest is collision-resistant, a computationally-bounded consensus adversary will not be able to convince a participant that a forged transcript is true.

Further, due to the forward-secrecy properties of our protocol, replay attacks will not be successful. Even messages with identical plaintexts will be encrypted with different ephemeral message keys, and therefore uniquely identifiable.

Privacy Adversary - The privacy adversary may participate in the conversation or control others in their participation. The goal of the privacy adversary is to convince a judge (either online - a current participant in the conversation, or offline - an entity that did not participate in the conversation, but has a complete transcript of the conversation available) that an honest participant in the conversation actually took part in the conversation by reading or sending messages to the other participants in the conversation.

In our protocol, all participants authenticate to each other via a triple ECDH mechanism¹⁶ using their long-term identity key, as well as an ephemeral ratchet key. Assuming the computational ECDH problem is hard, each participant is confident that all other members of the group have knowledge of their identity private keys and ephemeral ratchet keys (here proof of knowledge of the private identity key constitutes proof of identity). However, since the result of the authentication mechanism is the generation of a shared secret between two parties, it is possible for this authentication to be faked by one of the two parties. As a result, it is not possible to prove - in an information-theoretic sense - that any single party was authenticated at the key-agreement stage of the conversation. Thus, it is possible for a transcript of a conversation to be generated involving a particular participant that did not actually involve that participant. In other words, our protocol guarantees information-theoretic plausible deniability.

B. Network Model

The network model used by SMMP is that of a fully-connected graph. Operation of the network may be simplified somewhat by using a routing server to relay messages to all participants. The requirements for this routing server are not great. It should not have decryption capability, and as a result does not need to maintain its own ratchet state. It merely needs to route messages based on routing information supplied by the sending participant. Providing such routing information will not leak additional metadata to an adversary, because the adversary could also track the message flow along the connected graph, obtaining the same information.

The connected graph network suffers from scalability issues - for large enough groups, maintaining the ratchet state becomes very resource intensive. However, we do not believe this is a significant disadvantage, because person-to-person group conversations suffer from the same scalability problem. It would be unusual to see a group conversation between a number of participants larger than *e.g.* $N = 20$ or so. Larger conversations tend to break into sub-conversations. Extending SMMP to cover multiple sub-conversations with a group will be described in a later paper.

C. Cooperative Adversaries

Finally, we argue here that both Privacy and Consensus adversaries will cooperate with the SMMP protocol during the conversation. Although it is certainly possible for these entities to exhibit disruptive behavior (drop, duplicate, reorder, and replay messages as well as spam and DoS attacks), such disruptive actions would serve no purpose. The perpetrator of such acts will be immediately identifiable to the other members of the group, and faced with such attacks we expect that the group would immediately re-form without the attacking group member. Thus, the adversarial group member would lose the ability to follow the plaintext conversation. As a result all group members, adversarial or otherwise, are expected to faithfully follow the protocol.

XI. IMPLEMENTATION

Finally, we note that we have developed a reference implementation¹⁷ of SMMP. The reference implementation is built using the following six cryptographic primitives:

1. the AES256 symmetric encryption algorithm
2. ECDH on curve25519
3. DH with generator $g = 2$ where p is the 2048-bit MODP prime from RFC 3526
4. the SHA256 hash function
5. the HMAC message authentication code
6. the PBKDF2 key derivation function

Our reference implementation uses a simple TDMA algorithm for collision-avoidance with both conversation message and resynchronization packets.

* k0rx@uiowa.edu; Also at Department of Physics and Astronomy, The University of Iowa.

¹ J. Bian, R. Seker, and U. Topaloglu “Off-the-Record Instant Messaging for Group Conversation,” *IRI '07: Proceedings of Information Reuse and Integration*, pp. 79-84, IEEE Computer Society, 2007.

² I. Goldberg, M. D. Van Gundy, B. Ustanoglu, and H. Chen, “Multi-Party Off-the-Record Messaging,” *CSS '09: Proceedings of the 16th ACM Conference on Computer and Communication Security*, pp. 358-368, ACM, 2009.

³ Cryptocat Messaging Blog <https://github.com/cryptocat/mpotr>, accessed March, 2014.

⁴ H. Liu, E. Y. Vasserman, and N. Hopper, “Improved Group Off-the-Record Messaging,” *WPES'13: Proceedings of the ACM Workshop on Privacy in the Electronic Society*, ACM, 2013.

⁵ M. Burmester and Y. Desmedt, “A secure and efficient conference key distribution system,” *EUROCRYPT'94: Advances in Cryptology*, volume 950 of Lecture Notes in Computer Science, 1995.

- ⁶ “Pidgin,” <https://www.pidgin.im>, accessed March, 2014.
- ⁷ M. Blum, P. Feldman, and S. Micali, “Non-Interactive Zero-Knowledge and Its Applications,” *STOC 1988: Proceedings of the twentieth annual ACM symposium on Theory of computing*, pp. 103-112, ACM, 2008.
- ⁸ “Off the Record Messaging”, <https://otr.cyberpunks.ca/>, accessed March, 2014.
- ⁹ N. Borisov, I. Goldberg, and E. Brewer, “Off-the-Record Communication, or, Why Not To Use PGP,” *WPES’04: Proceedings of the ACM Workshop on Privacy in the Electronic Society*, ACM, 2004.
- ¹⁰ M. D. Raimondo, R. Gennaro, and H. Krawczyk, “Secure Off-the-Record Messaging,” *WPES’05 Proceedings of the 2005 ACM Workshop on Privacy in the Electronic Society*, pp. 81-89, ACM, 2005.
- ¹¹ C. Alexander and I. Goldberg, “Improved User Authentication in Off-The-Record Messaging,” *WPES’07: Proceedings of the ACM Workshop on Privacy in the Electronic Society*, ACM, 2007.
- ¹² R. Stedman, K. Yoshida, and I. Goldberg, “A User Study of Off-The-Record Messaging,” *SOUPS 2008: Symposium on Usable Privacy and Security*, pp. 1-10, SOUPS, 2008.
- ¹³ V. Moscaritolo, G. Belvin, and P. Zimmermann, “Silent Circle Instant Messaging Protocol Protocol Specification”, <https://silentcircle.com/static/download/SCIMP%20paper.pdf>, accessed March, 2014.
- ¹⁴ T. Perrin and M. Marlinspike, “Axolotl Ratchet”, <https://github.com/trevp/axolotl/wiki/newversion>, accessed March, 2014.
- ¹⁵ M. Just and S. Vaudenay, “Authenticated Multi-Party Key Agreement,” *ASIACRYPT’96 - Advances in Cryptology*, pp. 36-49 (1996).
- ¹⁶ T. Perrin and M. Marlinspike, “Simplifying OTR Deniability,” <https://whispersystems.org/blog/simplifying-otr-deniability/>, accessed March, 2014.
- ¹⁷ D. R. Andersen, M. S. Andersland, and T. J. Andersen, “Reference implementation of SMMP: A Secure Multi-party Messaging Protocol,” <https://notpublicyet>, 2014.

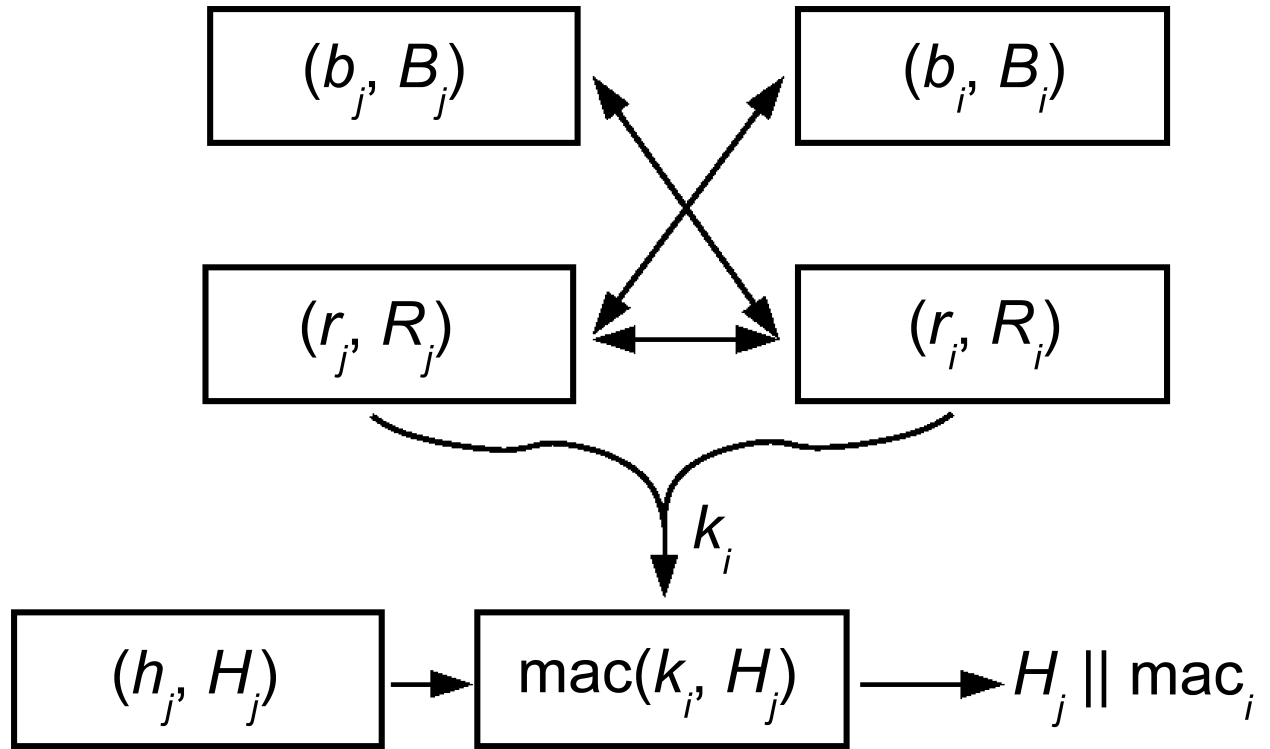


FIG. 1. Triple Diffie-Hellman authenticated key exchange.

Authenticated BD Group Key Exchange

Each participant generates: $(b_j, B_j), (r_j, R_j), (h_j, H_j)$




FIG. 2. Schematic illustration of the authenticated Burmester-Desmedt (BD) group key exchange algorithm for a group of three participants, A, B , and C . In the initial round, identity and ephemeral ratchet public keys are exchanged. A triple DH key is then computed for each participant, and the BD handshake public key mac is computed and attached to the BD handshake key, verifying its authenticity. The handshake public key, together with mac, is then exchanged in a second round. Using the public handshake keys, each participant computes his intermediate BD key. This intermediate key is then exchanged in a third round, and the group master key is computed by each user.